

# **Faster Transactions:**

## Query Tuning for Data Manipulation

Jeff Iannucci





# What is this session?

Improving performance for:

**INSERT**

**UPDATE**

**DELETE**























# *Le SQL Server*





Log Buffer Cache



Transaction Log



Buffer Cache




Pages on Disk



Locks






**Appetizers:** Alter the **target table**



**Entrees:** Manipulate the **transaction**



**Desserts:** Modify the **database settings**



**Appetizers:** Alter the target table



**Entrees:** Manipulate the transaction



**Desserts:** Modify the database settings

# Let's start with...

An **INSERT**

...of sorted data

...into an empty table

...that has no indexes

...(also known as a heap.)



# Clustered Index on a Heap

```
-- Create a table with a clustered index
CREATE TABLE TargetTable (
    ID int NOT NULL PRIMARY KEY CLUSTERED
    , Col1 varchar(100)
    , Col2 datetime
) ON [PRIMARY];

-- ...and then an ordered INSERT
INSERT SomeTable (Id, Col1, Col2)
SELECT ID, Col1, Col2
FROM SourceTable
ORDER BY ID;
```



# Clustered Index on a Heap

```
-- Create a table with a clustered index
CREATE TABLE TargetTable (
    ID int NOT NULL PRIMARY KEY CLUSTERED
    , Col1 varchar(100)
    , Col2 datetime
) ON [PRIMARY];

-- ...and then an ordered INSERT
INSERT SomeTable (Id, Col1, Col2)
SELECT ID, Col1, Col2
FROM SourceTable
ORDER BY ID;
```

# Clustered Index (CI) on Heap

Unsorted INSERTs are slower with a CI

Only faster when CI order matches sort order

Temporary tables can have a CI



# Now let's try a bit of...

An **UPDATE**

...to most or all rows of a table

...that has a clustered index

...and non-clustered indexes.



# Disabling indexes

```
-- Disable an index  
ALTER INDEX IX_SomeTable_ColX  
ON SomeSchema.SomeTable DISABLE;
```

```
UPDATE SomeTable  
SET ColX = 1  
WHERE ColX = 0;
```

```
-- Enable an index  
ALTER INDEX IX_SomeTable_ColX  
ON SomeSchema.SomeTable REBUILD;
```



# Disabling indexes

```
-- Disable an index  
ALTER INDEX IX_SomeTable_ColX  
ON SomeSchema.SomeTable DISABLE;
```

```
UPDATE SomeTable  
SET ColX = 1  
WHERE ColX = 0;
```

```
-- Enable an index  
ALTER INDEX IX_SomeTable_ColX  
ON SomeSchema.SomeTable REBUILD;
```

# Disabling indexes

REBUILDS take additional time

NCIs are not used until they are rebuilt

Disable all NCIs for INSERTs and DELETEs





**Appetizers:** Alter the target table



**Entrees:** Manipulate the transaction



**Desserts:** Modify the database settings

# Today let's have...

## A DELETE

...of a majority of data

...in a table with lots of rows

...with lots of user connections

...so don't slow other queries.





# Batch using WHILE with TOP

```
-- Batch using WHILE with TOP
```

```
DECLARE
```

```
    @BatchIdMin int = 1  
    , @BatchIdMax int  
    , @RowCount int = 1
```

```
WHILE (@RowCount > 0) BEGIN
```

```
    SELECT TOP (50000) @BatchIdMax = Id  
    FROM SomeTable  
    WHERE Id > @BatchIdMin  
    ORDER BY Id;
```

```
DELETE
```

```
FROM SomeTable
```

```
WHERE Id > @BatchIdMin  
      AND Id <= @BatchIdMax;
```

```
SET @RowCount = @@ROWCOUNT;
```

```
SET @BatchIdMin = @BatchIdMax
```

```
END;
```

# Batch using WHILE with TOP

```
-- Batch using WHILE with TOP
```

```
DECLARE
```

```
    @BatchIdMin int = 1  
    , @BatchIdMax int  
    , @RowCount int = 1
```

```
WHILE (@RowCount > 0) BEGIN
```

```
    SELECT TOP (50000) @BatchIdMax = Id  
    FROM SomeTable  
    WHERE Id > @BatchIdMin  
    ORDER BY Id;
```

```
DELETE
```

```
FROM SomeTable
```

```
WHERE Id > @BatchIdMin  
      AND Id <= @BatchIdMax;
```

```
SET @RowCount = @@ROWCOUNT;
```

```
SET @BatchIdMin = @BatchIdMax
```

```
END;
```

# Batch using WHILE with TOP

Will likely result in a slower query for you

Will certainly result in partial “transaction” if stopped

Experiment with different methods and batch sizes



# For RBAR dieters, try...

An **UPDATE**

...of much of the data

...in a table with lots of rows

...but each row is updated

...in a separate transaction.



# Explicit transaction

```
DECLARE @Id int = 1;

-- Explicit start
BEGIN TRANSACTION;

WHILE @id <= 100000 BEGIN

    EXEC usp_UpdateSomething @Id = @id;

    SET @id += 1;

END;

-- Explicit end
COMMIT;
```

# Explicit transaction

```
DECLARE @Id int = 1;

-- Explicit start
BEGIN TRANSACTION;

WHILE @id <= 100000 BEGIN

    EXEC usp_UpdateSomething @Id = @id;

    SET @id += 1;

END;

-- Explicit end
COMMIT;
```



# Explicit transaction

You are locking resources during your transaction

Concurrent queries will have to wait for resources

Could result in blocking or deadlocks



# What you really want is...

An **INSERT**

...of sorted or unsorted data

...into an empty table

...and make it fast as you can!




# Minimally logged INSERT

```
-- Minimal Logging with TABLOCK  
INSERT SomeTable  
  WITH (TABLOCK) (Col1, Col2, Col3)  
SELECT Col1, Col2, Col3  
FROM SomeOtherTable;
```

...“table lock”?

ROW  
LOCK



PAGE  
LOCK



TABLE  
LOCK



...“page allocation”?

RECORD logging (default)

Current LSN	Operation	Context
00000276:00008c68:00e7	LOP_INSERT_ROWS	LCX_HEAP
00000276:00008c68:00e8	LOP_INSERT_ROWS	LCX_HEAP
00000276:00008c68:00e9	LOP_INSERT_ROWS	LCX_HEAP
00000276:00008c68:00ea	LOP_INSERT_ROWS	LCX_HEAP
00000276:00008c68:00eb	LOP_INSERT_ROWS	LCX_HEAP
00000276:00008c68:00ec	LOP_INSERT_ROWS	LCX_HEAP
00000276:00008c68:00ed	LOP_INSERT_ROWS	LCX_HEAP
00000276:00008c68:00ee	LOP_INSERT_ROWS	LCX_HEAP
00000276:00008c68:00ef	LOP_INSERT_ROWS	LCX_HEAP

PAGE allocation logging

Current LSN	Operation	Context
00000276:00008ab8:004d	LOP_MODIFY_ROW	LCX_PFS
00000276:00008ab8:004e	LOP_SET_BITS	LCX_IAM
00000276:00008ab8:004f	LOP_SET_BITS	LCX_GAM

# A few minimally-logged transaction rules

Database can NOT be in FULL recovery model

Table is not replicated

Table is not memory-optimized

Table has no indexes (is a heap), or...

...if the table has indexes, then it must be empty

May require Trace Flag 610

<https://learn.microsoft.com/en-us/sql/relational-databases/import-export/prerequisites-for-minimal-logging-in-bulk-import?view=sql-server-ver16>



# Minimally logged INSERT

```
-- Minimal Logging with TABLOCK  
INSERT SomeTable  
WITH (TABLOCK) (Col1, Col2, Col3)  
SELECT Col1, Col2, Col3  
FROM SomeOtherTable;
```

# Minimally logged INSERT

```
-- Minimal Logging with TABLOCK  
INSERT SomeTable  
WITH (TABLOCK) (Col1, Col2, Col3)  
SELECT Col1, Col2, Col3  
FROM SomeOtherTable;
```

# Minimally logged INSERT

No guarantees of minimal logging with all those rules

Concurrent queries will likely have to wait to use the table

Minimal logging only works for INSERTs



# Next you might enjoy...

A **DELETE**

...of a majority of data

...in a table with lots of rows

...and make it fast

...like minimal logging!



# Minimally logged DELETE

```
-- 1st INSERT WITH (TABLOCK)
```

```
INSERT NewTable  
WITH (TABLOCK) (ID, Col1)  
SELECT ID, Col1  
FROM SomeTable  
WHERE Col1 = 'X';
```

```
-- TRUNCATE
```

```
TRUNCATE TABLE SomeTable;
```

```
-- 2nd INSERT WITH (TABLOCK)
```

```
SET IDENTITY_INSERT SomeTable ON;
```

```
INSERT SomeTable  
WITH (TABLOCK) (ID, Col1)  
SELECT ID, Col1  
FROM NewTable;
```

```
SET IDENTITY_INSERT SomeTable OFF;
```



# Minimally logged DELETE

```
-- 1st INSERT WITH (TABLOCK)
```

```
INSERT NewTable  
WITH (TABLOCK) (ID, Col1)  
SELECT ID, Col1  
FROM SomeTable  
WHERE Col1 = 'X';
```

```
-- TRUNCATE
```

```
TRUNCATE TABLE SomeTable;
```

```
-- 2nd INSERT WITH (TABLOCK)
```

```
SET IDENTITY_INSERT SomeTable ON;
```

```
INSERT SomeTable  
WITH (TABLOCK) (ID, Col1)  
SELECT ID, Col1  
FROM NewTable;
```

```
SET IDENTITY_INSERT SomeTable OFF;
```

# Minimally logged DELETE

All the rules for minimal logging apply

Best for removing a majority of the data

DISABLE/re-ENABLE constraints





**Appetizers:** Alter the target table



**Entrees:** Manipulate the transaction



**Desserts:** Modify the database settings

# What sounds good is...

An **INSERT**

...of sorted or unsorted data

...into an empty table

...in a database

...with FULL recovery model

...and give me minimal logging!



# Bulk Logged Recovery Model

```
-- Change recovery model
```

```
ALTER DATABASE SomeDB  
SET RECOVERY BULK_LOGGED;
```

```
INSERT SomeTable  
WITH (TABLOCK) (Col1, Col2, Col3)  
SELECT Col1, Col2, Col3  
FROM SomeOtherTable;
```

```
ALTER DATABASE SomeDB  
SET RECOVERY FULL;
```



# Bulk Logged Recovery Model

```
-- Change recovery model
```

```
ALTER DATABASE SomeDB  
SET RECOVERY BULK_LOGGED;
```

```
INSERT SomeTable  
WITH (TABLOCK) (Col1, Col2, Col3)  
SELECT Col1, Col2, Col3  
FROM SomeOtherTable;
```

```
ALTER DATABASE SomeDB  
SET RECOVERY FULL;
```

# Bulk Logged Recovery Model

ALL database transactions  
affected

No Point-In-Time recovery  
while in Bulk Logged

Consult with your Database  
Administrator



# Let's boldly finish with...

A **DELETE**

...of a bunch of data

...in a table with lots of rows

...and then issue a **ROLLBACK**.



# Accelerated Database Recovery (ADR)

```
ALTER DATABASE SomeDB  
SET ACCELERATED_DATABASE_RECOVERY = ON;
```

```
-- Start a transaction  
BEGIN TRANSACTION
```

```
DELETE  
FROM SomeTable  
WHERE Id = @SomeVariable
```

```
-- ROLLBACK  
ROLLBACK
```

# Accelerated Database Recovery (ADR)

```
ALTER DATABASE SomeDB  
SET ACCELERATED_DATABASE_RECOVERY = ON;
```

```
-- Start a transaction  
BEGIN TRANSACTION
```

```
DELETE  
FROM SomeTable  
WHERE Id = @SomeVariable
```

```
-- ROLLBACK  
ROLLBACK
```

# Accelerated Database Recovery

Writing more data during your transaction

Requires more space in data files

Database level setting





# Menu

## Antipasti

Platto di carne italiana  
verduro grigliato all'aceto balsamico  
Mozzarella di Bufala e pesto  
Bruschetta

## Secondi

Pollo alla griglia  
Pollo brasato  
Rizotto ai funghi e ricotta  
Spagetti al forno  
Pasta di semola di grano duro alla contadina





## Dolci

Crema di fragole  
Torta di limone  
Fruite fresca











# Appetizers: Alter the target table

## Query options

	INSERT	UPDATE	DELETE
Clustered index on a heap			
Disable indexes			






# Entrees: Transaction manipulation

## Query options

	INSERT	UPDATE	DELETE
Batching (WHILE with TOP)			
Explicit transactions			
Minimally logged INSERT			
Minimally logged DELETE			

# Desserts: Database settings

## Query options

	INSERT	UPDATE	DELETE
Bulk Logged recovery model			
Accelerated Database Recovery			



# Transaction cost multipliers

Triggers

Replication

Change Data Capture

Availability Groups







[desertrdb.com](http://desertrdb.com)



[jeff@desertrdb.com](mailto:jeff@desertrdb.com)



[@desertrdb](https://twitter.com/desertrdb)

