

Making the Most of Query Store in the Real World

Jeff Iannucci

Who in the world is Jeff Iannucci?




Consultant at Straight Path Solutions



Content Author at Pluralsight



Author of “Learn SQL in a Month of Lunches”



github.com/desertdba



In the “**Presentations**” folder:

- All the slides!
- One handy T-SQL script!

A word of thanks to these folks



Conor Cunningham



Erik Darling



Grant Fritchey



Tara Kizer



Kendra Little



Brent Ozar



Paul Randal



Erin Stellato

Has this been your experience?

- Query Store made performance worse
- Query Store unexpectedly went READ_ONLY
- Forced query plans didn't always get used
- Can we solve any REAL problems with this thing?



YES, WE CAN!!!

“Real world” solutions

- Ways to use Query Store beyond forcing plans
- Considerations for defaults and “best practices”
- Queries and free, community-supported tools
- Common problems you can solve

Setting up Query Store



What defaults to **change**...maybe

- MAX_STORAGE_SIZE_MB
- MAX_PLANS_PER_QUERY
- STALE_QUERY_THRESHOLD_DAYS
- QUERY_CAPTURE_MODE

What defaults to **leave alone**

- DATA_FLUSH_INTERVAL_SECONDS
- INTERVAL_LENGTH_MINUTES
- SIZE_BASED_CLEANUP_MODE
- WAIT_STATS_CAPTURE_MODE

Trace flags

- Query Store uses in-memory tables
- **TF 7752** – allows queries to execute while QS loads into memory (SQL Server 2016 & 2017 only)
- **TF 7745** – bypasses writing to disk at shutdown, losing unflushed data (all versions)

dba-tools (PowerShell)

- `Get-DbadbQueryStoreOption`
- `Set-DbadbQueryStoreOption`
- `Copy-DbadbQueryStoreOption`
- `Test-DbadbQueryStore`

You should get a job (for monitoring)

```
/*  
Query Store - current status  
*/  
SELECT  
    DB_NAME() as DatabaseName  
    , actual_state_desc  
    , readonly_reason  
    , max_storage_size_mb  
    , (max_storage_size_mb - current_storage_size_mb) AS query_store_free_space_mb  
    , flush_interval_seconds  
    , interval_length_minutes  
    , stale_query_threshold_days  
    , max_plans_per_query  
    , query_capture_mode_desc  
    , size_based_cleanup_mode_desc  
FROM sys.database_query_store_options
```



Tools you can use

SSMS built-in reports

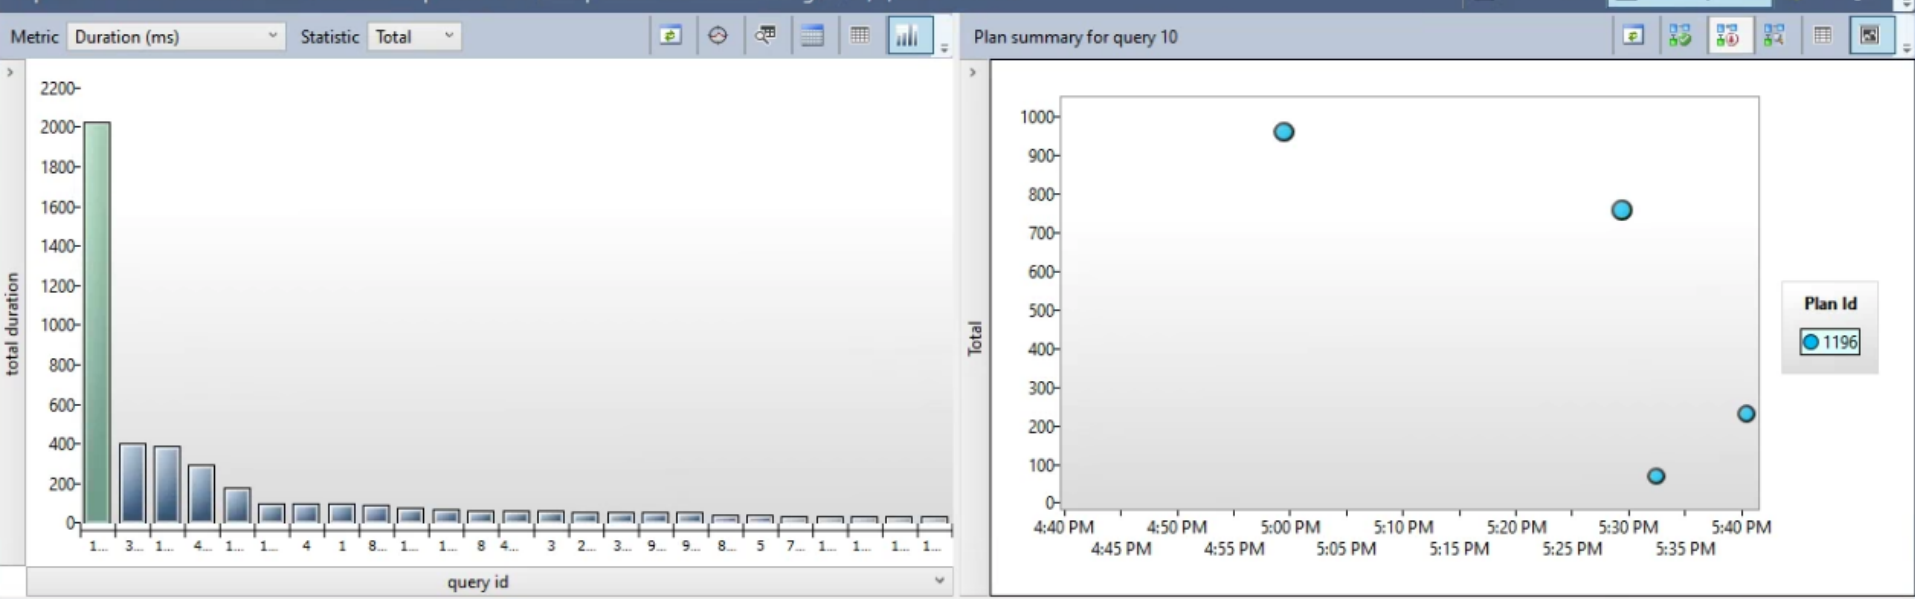
- Top Resource Consuming Queries
- Queries With Forced Plans
- Tracked Queries

SSMS built-in reports

- Top Resource Consuming Queries
- Queries With Forced Plans
- Tracked Queries

Top Resource Consuming Queries

Top 25 resource consumers for database ReportServer. Time period: Last hour ending at 11/1/2023 5:40 PM



Defaults are...

- Top 25
- Last hour
- Total and Duration (2 separate defaults)
- Cute and colorful circles (squares and triangles too)

What is the real problem?

- CPU usage
- Memory consumption
- I/O (reads and/or writes)
- Total, Average, Maximum

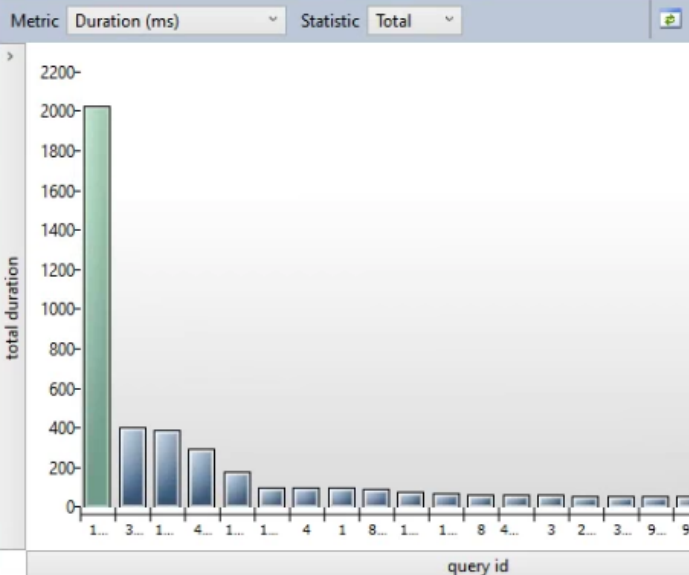
Paul Randal's Wait Stats

- <https://www.sqlskills.com/blogs/paul/wait-statistics-or-please-tell-me-where-it-hurts/>

Wait Type	Wait_S	Resource_S	Signal_S	WaitCount	Percentage	AvgWait_S	AvgRes_S	AvgSig_S	Help/Info URL
CXPACKET	7668440.60	6793057.13	875383.47	2362131893	65.03	0.0032	0.0029	0.0004	https://www.sqlskills.com/help/waits/CXPACKET
SOS_SCHEDULER_YIELD	2127990.30	1678.10	2126312.20	2777370420	18.05	0.0008	0.0000	0.0008	https://www.sqlskills.com/help/waits/SOS_SCHEDULER_YIELD
PAGEIOLATCH_SH	449569.26	421665.76	27903.51	460398999	3.81	0.0010	0.0009	0.0001	https://www.sqlskills.com/help/waits/PAGEIOLATCH_SH
ASYNC_NETWORK_IO	351253.40	328880.93	22372.48	160290986	2.98	0.0022	0.0021	0.0001	https://www.sqlskills.com/help/waits/ASYNC_NETWORK_IO
MSQL_XP	180808.78	180808.78	0.00	4249985	1.53	0.0425	0.0425	0.0000	https://www.sqlskills.com/help/waits/MSQL_XP
WRITELOG	180553.20	130858.81	49694.39	229807177	1.53	0.0008	0.0006	0.0002	https://www.sqlskills.com/help/waits/WRITELOG
OLEDB	137839.73	137839.73	0.00	10707789161	1.17	0.0000	0.0000	0.0000	https://www.sqlskills.com/help/waits/OLEDB
LATCH_EX	94723.08	68904.50	25818.58	316396002	0.80	0.0003	0.0002	0.0001	https://www.sqlskills.com/help/waits/LATCH_EX
PAGEIOLATCH_EX	85873.68	84007.76	1865.92	137039849	0.73	0.0006	0.0006	0.0000	https://www.sqlskills.com/help/waits/PAGEIOLATCH_EX

Top Resource Consumers

Top 25 resource consumers for database ReportServer. Time period: Last



Configure Top Resource Consumption Criteria

Resource Consumption Criteria

Check for top consumers of:

- ☐ Execution Count
- ☒ Duration (ms)
- ☐ CPU Time (ms)
- ☐ Logical Reads (KB)
- ☐ Logical Writes (KB)
- ☐ Physical Reads (KB)
- ☐ CLR Time (ms)
- ☐ DOP
- ☐ Memory Consumption (KB)
- ☐ Row Count

Based on:

- ☐ Avg
- ☐ Max
- ☐ Min
- ☐ Std Dev
- ☒ Total

Time Interval

Last hour | From: | To: |

Time Format: ☒ Local ☐ UTC

Return

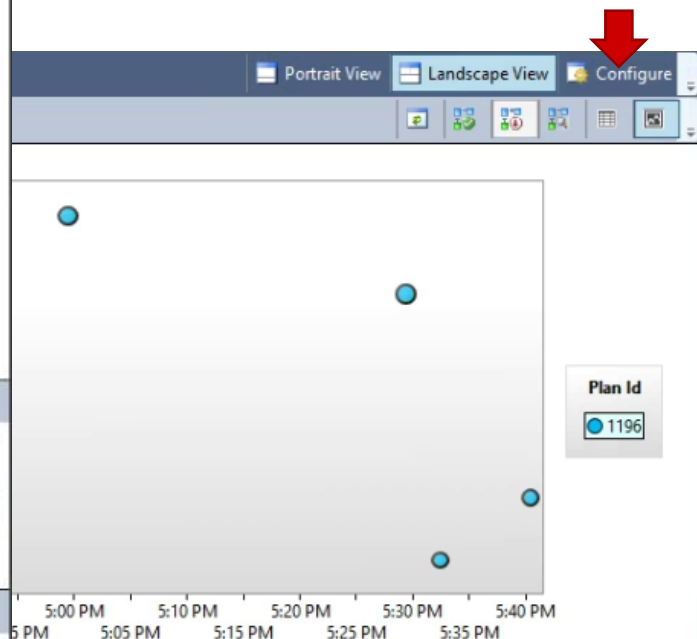
☐ All

☒ Top 25

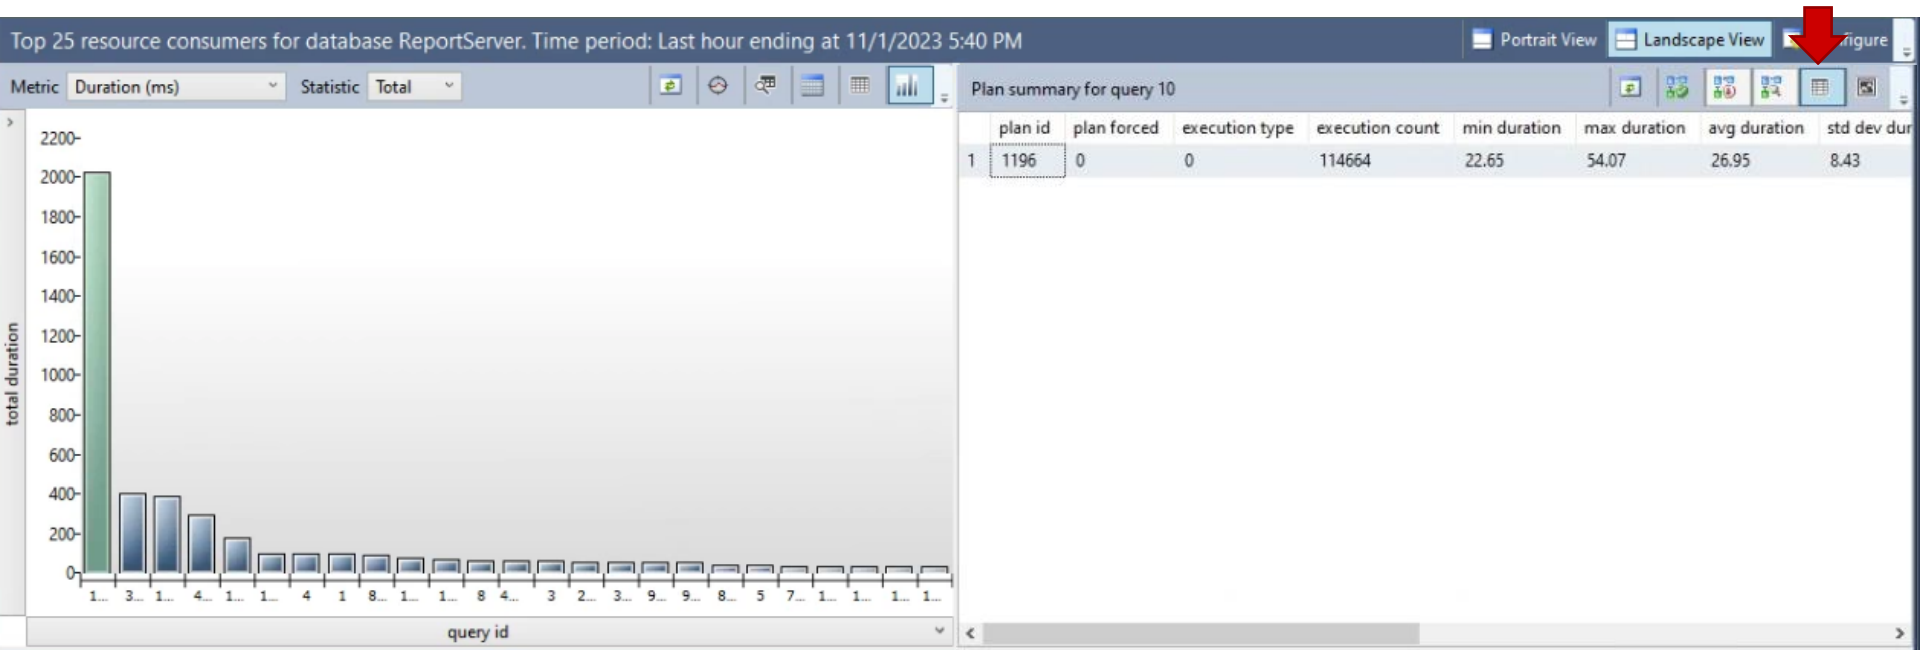
Filters

Minimum number of query plans: 1

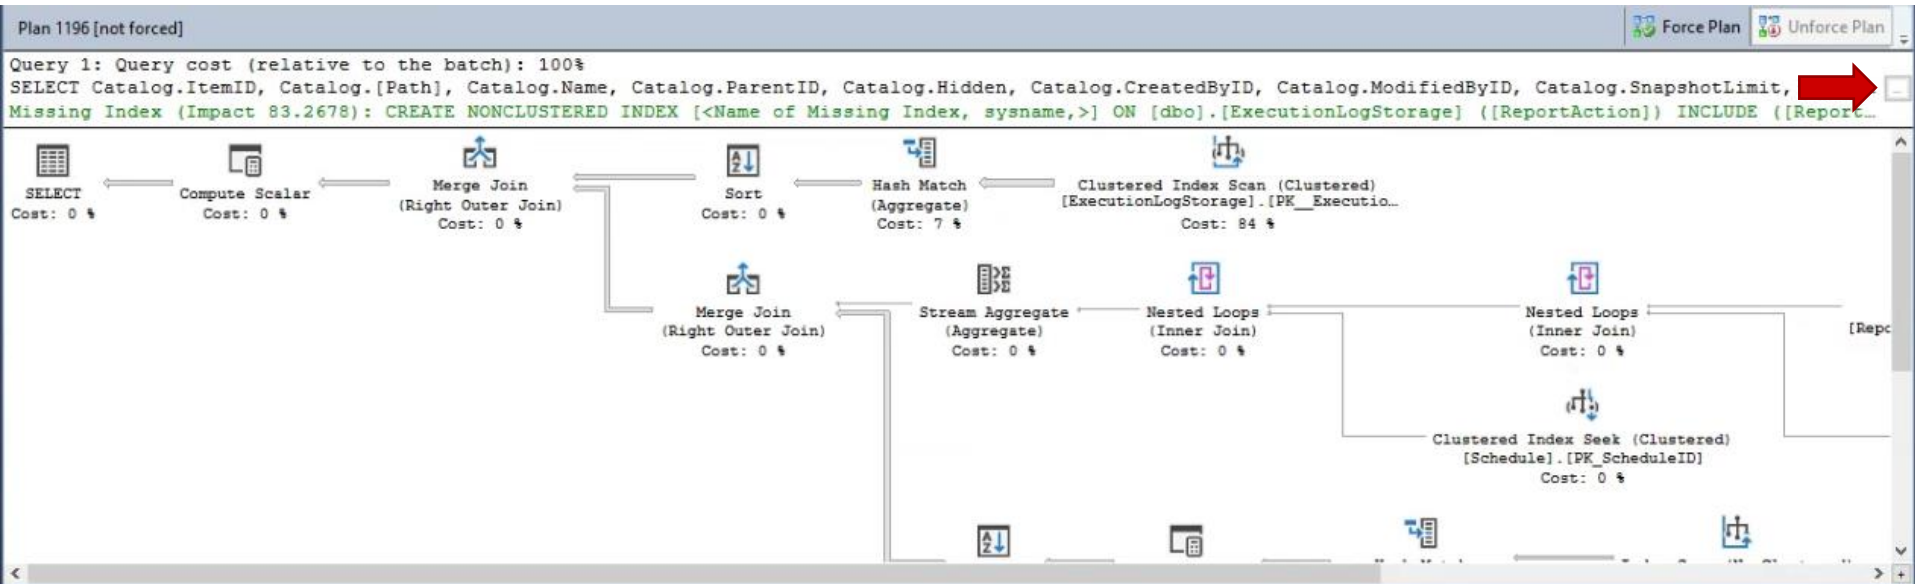
Ok Cancel Apply



Top Resource Consuming Queries



Top Resource Consuming Queries

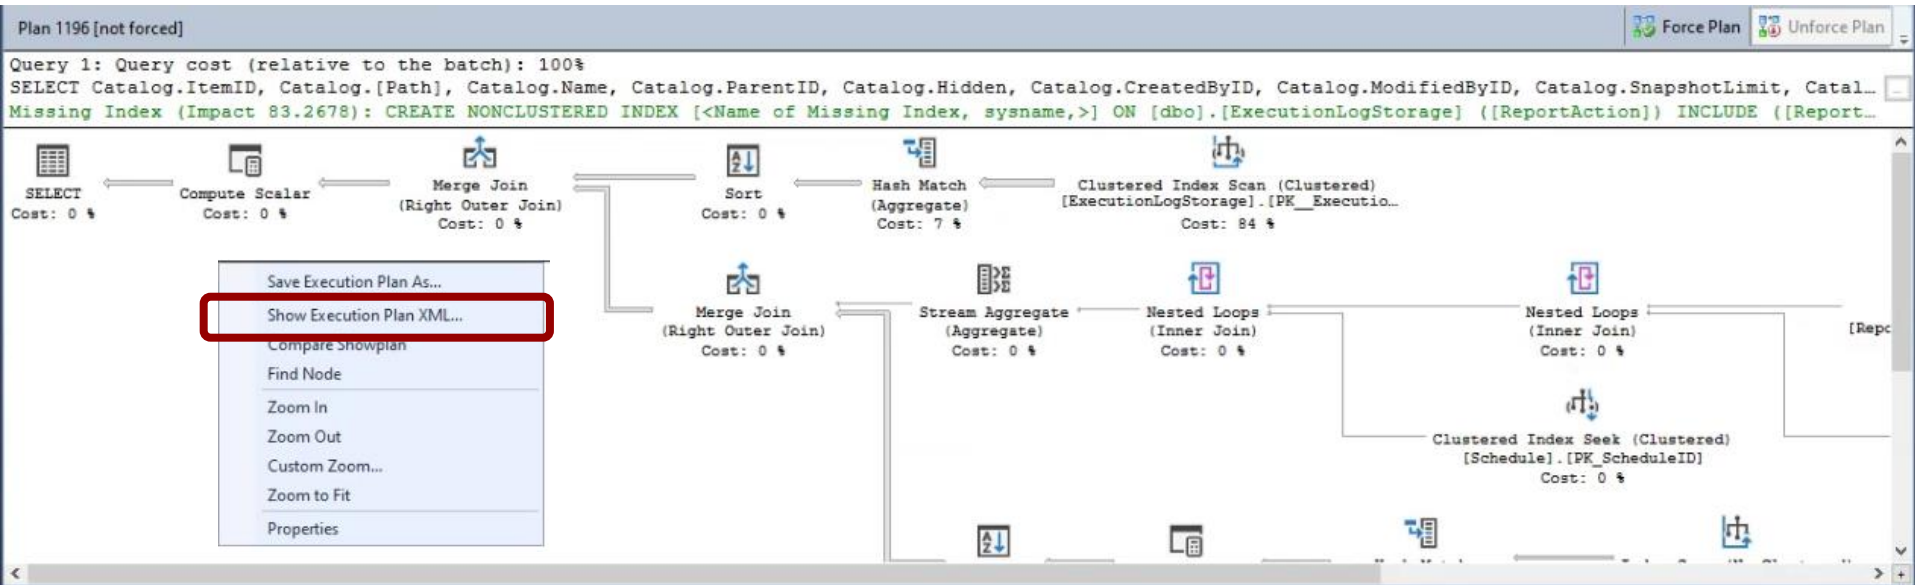


Top Resource Consuming Queries

```
/*  
This query text was retrieved from showplan XML, and may be truncated.  
*/
```

```
SELECT  
  
Catalog.ItemID,  
Catalog.[Path],  
Catalog.Name,  
Catalog.ParentID,  
Catalog.Hidden,  
Catalog.CreatedByID,  
Catalog.ModifiedByID,  
Catalog.SnapshotLimit,  
Catalog.PolicyID,  
Catalog.PolicyRoot,  
Catalog.ExecutionFlag,  
Catalog.ExecutionTime,  
Catalog.CreationDate,  
Catalog.ModifiedDate,  
Catalog.[Description],  
(SELECT UserName FROM Users WHERE UserID = Catalog.CreatedByID) CreatedByName,  
(SELECT UserName FROM Users WHERE UserID = Catalog.ModifiedByID) ModifiedByName,  
(SELECT MAX(NextRunTime) FROM Schedule INNER JOIN ReportSchedule ON  
Schedule.ScheduleID = ReportSchedule.ScheduleID  
WHERE ReportSchedule.ReportID = Catalog.ItemID) NextRunTime,  
(SELECT MAX(TimeStart) FROM ExecutionLog WHERE ExecutionLog.ReportID = Catalog.ItemID) LastExecutionTime  
FROM Catalog Catalog WITH (NOLOCK) WHERE Catalog.Type = 4 OR Catalog.Type = 2
```


Top Resource Consuming Queries



Top Resource Consuming Queries

WARNING: This
could contain PII

```
<Identifier>
  <ColumnReference Column="@AuthType" />
</Identifier>
</ScalarOperator>
</RangeExpressions>
</Prefix>
</SeekKeys>
</SeekPredicateNew>
</SeekPredicates>
</IndexScan>
</RelOp>
</NestedLoops>
</RelOp>
</ComputeScalar>
</RelOp>
<ParameterList>
  <ColumnReference Column="@AuthType" ParameterDataType="int" ParameterCompiledValue="(1)" />
  <ColumnReference Column="@EditSessionID" ParameterDataType="varchar(32)" ParameterCompiledValue="NULL" />
  <ColumnReference Column="@Path" ParameterDataType="nvarchar(425)" ParameterCompiledValue="N'/CustomerReports/Medallion/ElkayDailyOrderListing'" />
</ParameterList>
</QueryPlan>
</StmtSimple>
</Statements>
</Batch>
</BatchSequence>
</ShowPlanXML>
```

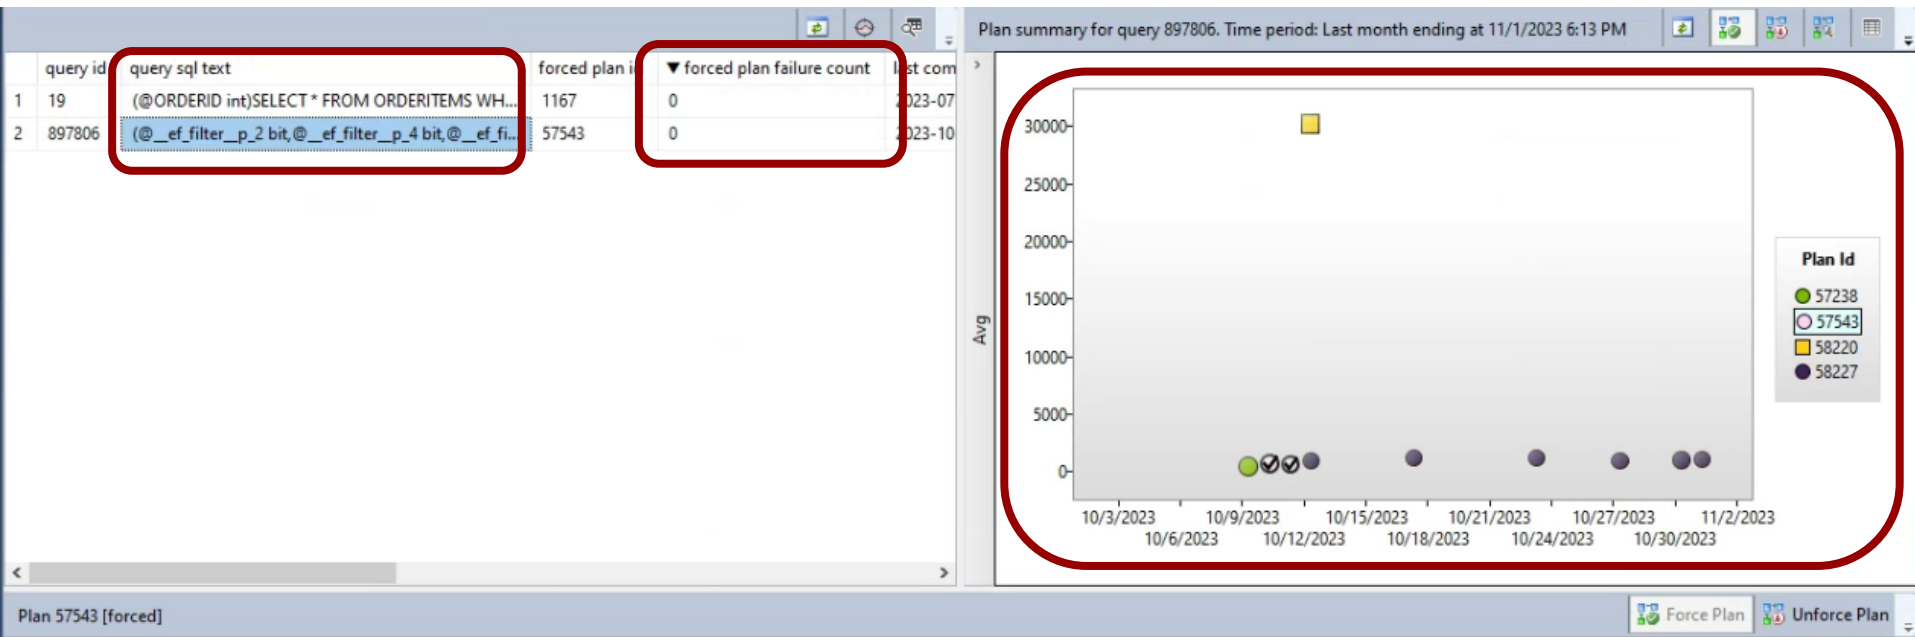
SSMS built-in reports

- Top Resource Consuming Queries
- Queries With Forced Plans
- Tracked Queries

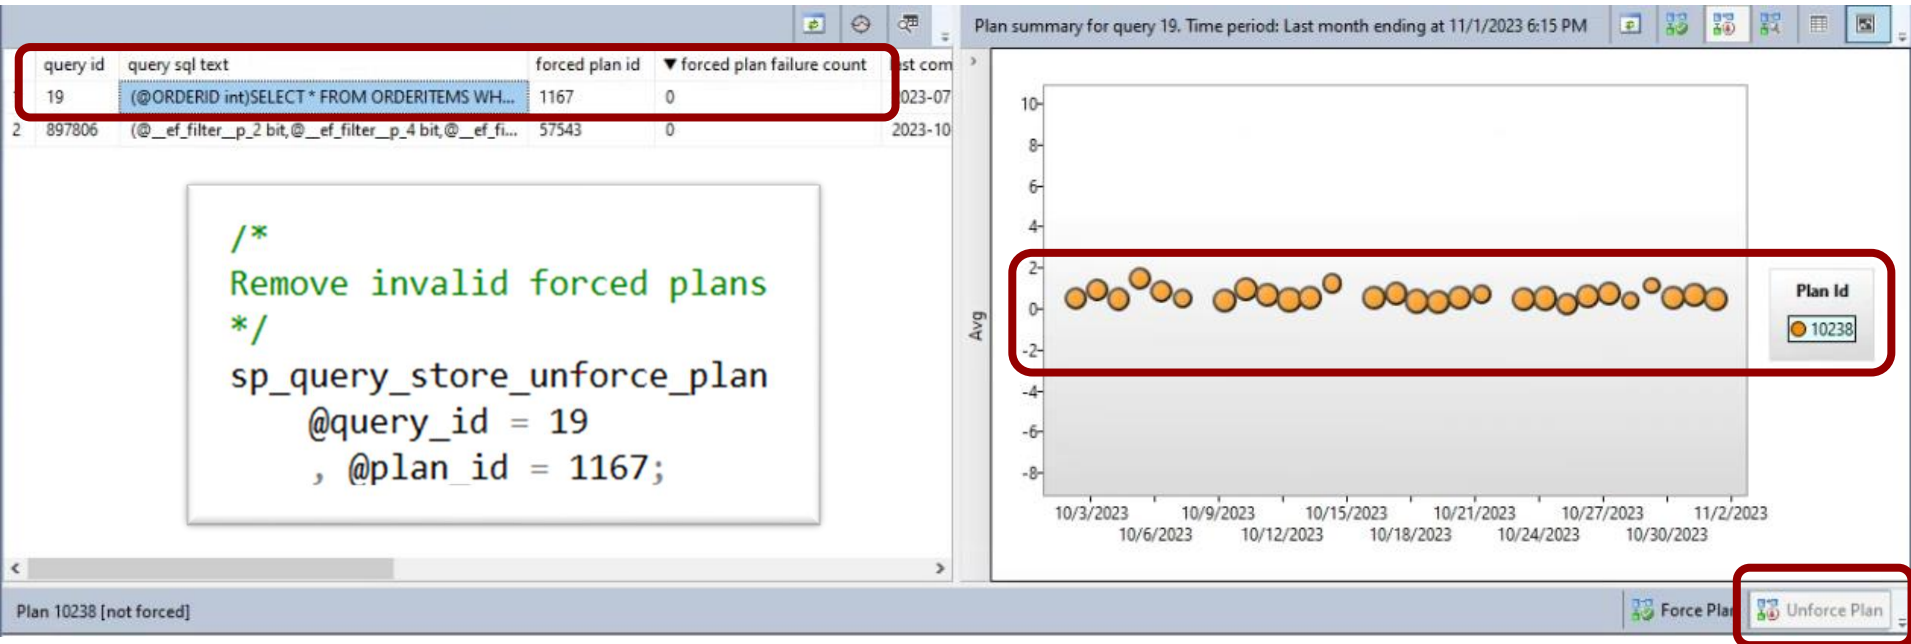


**Forced plans
are rarely a
long-term
solution**

Queries With Forced Plans



Queries With Forced Plans



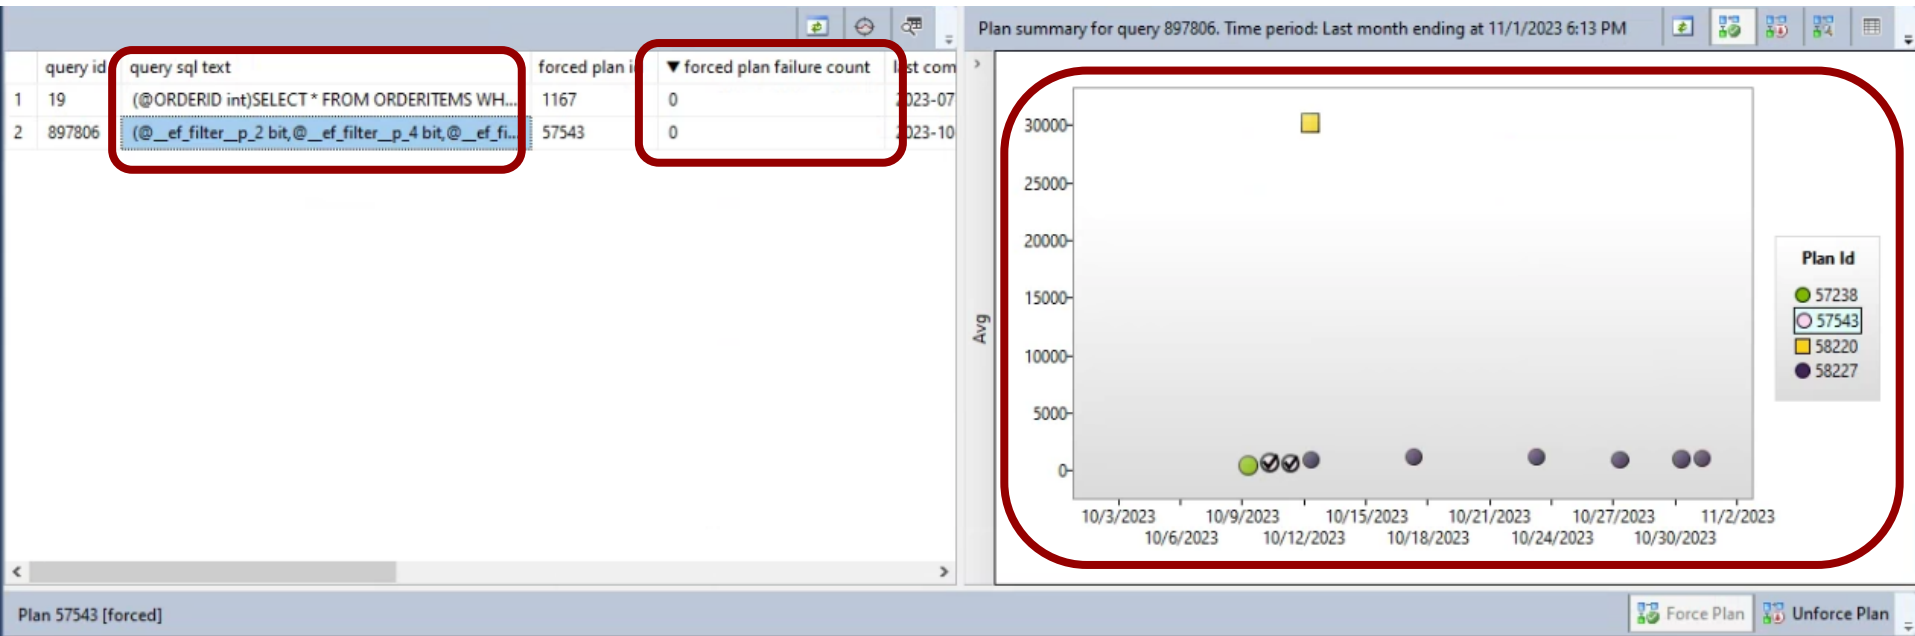
Failed forced plans

```
/*  
Find failed forced plans  
*/  
SELECT  
    p.query_id  
    , p.plan_id  
    , q.object_id as containing_object_id  
    , p.force_failure_count  
    , p.last_force_failure_reason  
    , p.last_force_failure_reason_desc  
    , p.last_execution_time  
FROM sys.query_store_plan AS p  
JOIN sys.query_store_query AS q  
    ON p.query_id = q.query_id  
WHERE p.is_forced_plan = 1  
      AND p.force_failure_count > 0;
```

Failed forced plans

Results		Messages					
	query_id	plan_id	containing_object_id	force_failure_count	last_force_failure_reason	last_force_failure_reason_desc	last_execution_time
1	94515	41042	0	1746	8698	NO_PLAN	2023-01-06 21:10:21.9730000 +00:00
2	91452	55877	0	204	8712	NO_INDEX	2023-02-09 20:11:52.1170000 +00:00
3	5149	57907	0	1349	8712	NO_INDEX	2023-02-21 16:07:42.7870000 +00:00
4	5410	91409	0	593	8698	NO_PLAN	2023-02-16 19:22:34.4100000 +00:00
5	12	9117995	1216540659	10	8695	GENERAL_FAILURE	2023-07-27 15:57:14.7330000 +00:00
6	17	9118009	1216540659	6	8695	GENERAL_FAILURE	2023-07-27 15:52:17.2800000 +00:00
7	23908674	9236540	1216540659	2184	8695	GENERAL_FAILURE	2023-10-28 16:52:10.7700000 +00:00
8	23908673	9534016	1216540659	287	8695	GENERAL_FAILURE	2023-10-30 12:07:16.6700000 +00:00
9	23908676	9534028	1216540659	277	8695	GENERAL_FAILURE	2023-10-30 12:07:15.9100000 +00:00

But...what's the deal here?





“Morally equivalent plans”








- Query Store has some flexibility
- Not a bug, but a feature
- Don't show as failures
- Forced plans are rarely a long-term solution

SSMS built-in reports

- Top Resource Consuming Queries
- Queries With Forced Plans
- Tracked Queries

Tracked Queries

Tracking query  

 Auto-Refresh  Refresh  Force Plan  Unforce Plan  Compare Plans  View Query  Configure

Tracked Queries

- Shows performance of one specific `query_id`
- Can compare plans
- Use “Auto-Update” to track in real time

Find queries (**query_id**) by...

- Object name (stored procedure/function/trigger)
- String of characters
- Time of execution

...by object name

```
/*  
Find query_id by object name (stored procedure, function, trigger, etc.)  
*/  
SELECT q.query_id  
       , t.query_sql_text  
FROM sys.query_store_query AS q  
JOIN sys.query_store_query_text AS t  
      ON q.query_text_id = t.query_text_id  
WHERE q.object_id = OBJECT_ID('zzz');
```

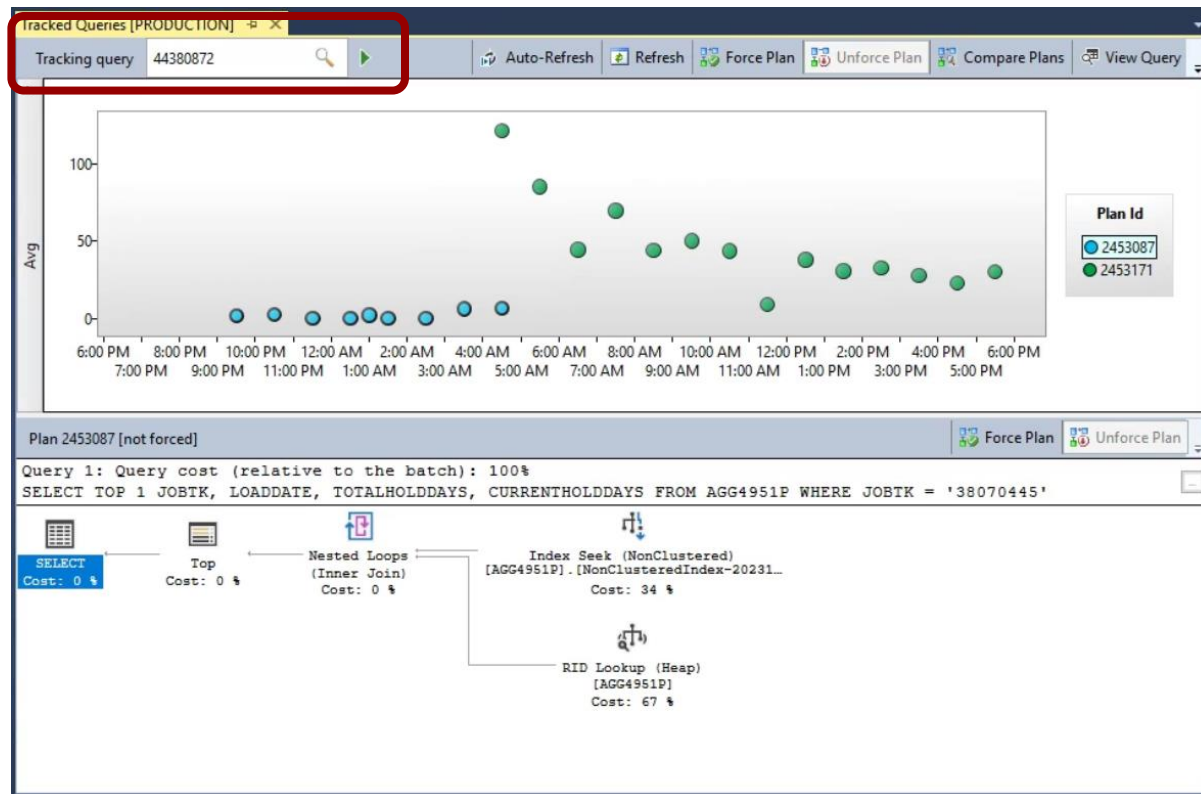
...by string

```
/*  
Find query_id by string  
*/  
SELECT q.query_id  
       , t.query_sql_text  
FROM sys.query_store_query AS q  
JOIN sys.query_store_query_text AS t  
      ON q.query_text_id = t.query_text_id  
WHERE t.query_sql_text LIKE '%zzz%';
```


...by time of execution

```
/*  
Find query_id by time of execution  
*/  
SELECT q.query_id  
      , t.query_sql_text  
FROM sys.query_store_query AS q  
JOIN sys.query_store_query_text AS t  
      ON q.query_text_id = t.query_text_id  
WHERE q.last_execution_time BETWEEN '1900/01/30 23:00:00' AND '1900/01/31 01:00:00';
```

Tracked Queries



Tracked Queries - Comparison

Showplan Comparison - Tracked Queries [PRODUCTION]

Plan 2453087
SELECT TOP 1 JOBTK, LOADDATE, TOTALHOLDDAYS, CURRENTHOLDDAYS FROM AGG4951P WH...

SELECT Cost: ...

Top Cost: ...

Nested L... (Inner J... Cost: 0 %

Index Seek (NonClustered) [AGG4951P].[NonClusteredIndex-20... Cost: 34 %

RID Lookup... [AGG4951P] Cost: 67 %

Plan 2453171
SELECT TOP 1 JOBTK, LOADDATE, TOTALHOLDDAYS, CURRENTHOLDDAYS FROM AGG4951P WH...

SELECT Cost: ...

Top Cost: ...

Table S... [AGG495... Cost: 1...

Properties

Top Plan
SELECT

Bottom Plan
SELECT

Actual 0
Cach: 32 KB
Card: 150
Comp: 4
Comp: 312
Comp: 6
Data: 1
Estim: 0
Estim: 1
Estim: 0 (0%)
Estim: 0.00973

Memc
Optim FULL
Optim
Optim
Param: 0
Query: 0x5F51A4182
Query: 0x7B036
Reas: Good Enough
Retri: false
Secu: False
Set O ANSI_NULLS
State: SELECT TOP
State: 0
State: 0x090058BA4
Trace

Actual Number of
Actual number of rows
for All Executions
output by this operator.
For rows of type
PLAN_ROWS only.

Actual 0
Cach: 24 KB
Card: 150
Comp: 2
Comp: 304
Comp: 6
Data: 1
Estim: 0
Estim: 1
Estim: 0 (0%)
Estim: 0.00794

Mem
Optim FULL
Optim
Optim
Param: 0
Query: 0x5F51A4182
Query: 0xA687
Reas: Good Enough
Retri: false
Secu: False
Set C ANSI_NULLS
State: SELECT TOP
State: 0
State: 0x090058BA4
Trace

Actual Number of
Actual number of rows
for All Executions
output by this
operator. For rows
of type PLAN_ROWS ...

sp_BlitzQueryStore

- Prioritized listing of potential issues
- Grouped by findings (“Worst Avg CPU”, etc..)
- Better if you don’t know where to start
- <https://github.com/BrentOzarULTD/SQL-Server-First-Responder-Kit>

sp_QuickieStore

- Choose a metric (CPU, reads, writes, memory...)
- Returns query text by default
- Faster if you know what you're looking for
- https://github.com/erikdarlingdata/DarlingData/tree/main/sp_QuickieStore



**Common
problems
you can
solve**

Testing index changes

- Missing index requests shown in plans
- Review existing indexes before adding
- Unforce plans if adding/modifying indexes

Reducing upgrade pains

- Upgrade to higher version of SQL Server
- ...but don't raise the Compatibility Level
- If performance is stable, raise the compatibility level
- Use Query Store for analysis and regression fixes

Deadlocks

- You can capture deadlocks with extended events
- Query Store can give you more info
- Use `sys.query_store_runtime_stats`
- WHERE `execution_type = 4`

What do those shapes mean?

- Successful query execution (0)
- Aborted query (3)
- ▲ Error during execution (4)

Deadlocks

```
/*  
Find deadlocks (and other aborted/cancelled queries)  
*/  
SELECT  
    q.query_id  
    , t.query_sql_text  
    , r.execution_type  
    , r.execution_type_desc  
    , x.query_plan_xml  
    , r.count_executions  
    , r.last_execution_time  
FROM sys.query_store_query q  
JOIN sys.query_store_plan p  
    ON q.query_id=p.query_id  
JOIN sys.query_store_query_text t  
    ON q.query_text_id=t.query_text_id  
OUTER APPLY (SELECT TRY_CONVERT(XML, p.query_plan) AS query_plan_xml) x  
JOIN sys.query_store_runtime_stats r  
    ON p.plan_id = r.plan_id  
WHERE r.execution_type = 4 /* Exception aborted execution */  
    AND q.last_execution_time > GETDATE() - 1;
```

Deadlocks

- **WARNING:** this data is still aggregated
- Compare deadlocks using **sql_text**
- Review execution plans for objects locked in deadlocks
- Use **Tracked Queries** report

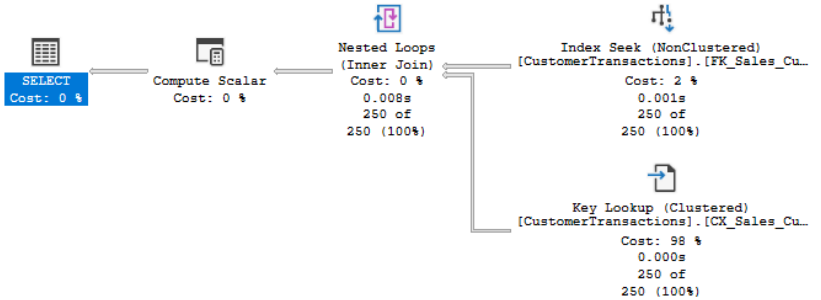
Parameter Sniffing

EXEC GetTransactionByCustomer

@CustomerID = 401

Query 1: Query cost (relative to the batch): 100%

SELECT * FROM [Sales].[CustomerTransactions] WHERE CustomerID = @CustomerID




EXEC GetTransactionByCustomer

@CustomerID = 976

Query 1: Query cost (relative to the batch): 100%

SELECT * FROM [Sales].[CustomerTransactions] WHERE CustomerID = @CustomerID
Missing Index (Impact 99.3199): CREATE NONCLUSTERED INDEX [<Name of Missing Index,



- Different plans, different performance

What about FORCE_LAST_GOOD_PLAN?

- Automatically force a better plan when found
- Can still choose a “bad” plan
- Can take hours to automatically revert

Parameter sniffing

- Query `sys.dm_db_tuning_recommendations`
- Insert JSON data from Details column into a table
- Create a job that queries the table
- Set job to “fail” when regressions occur

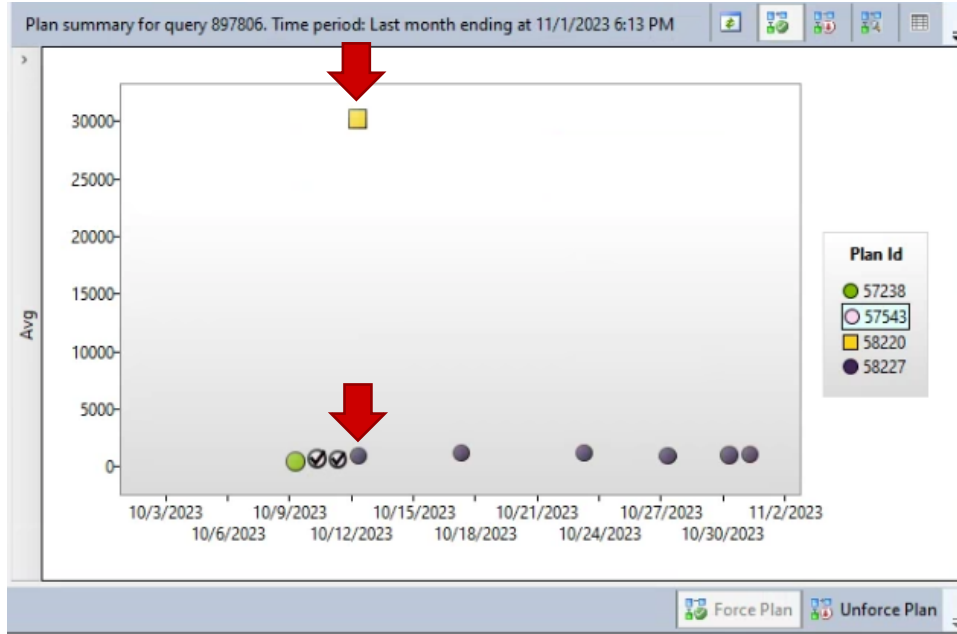
vor Energy VDI

```


11.
12. DECLARE @current_date datetime = GETDATE(), @cmd nvarchar(max), @query nvarchar(1000), @database
    sysname;
13.
14. SET @database = 'DB_Administration'
15. SET @cmd = '
16.
17. INSERT INTO ' + @database + '.dbo.AutomaticTuning
18. SELECT *
19. FROM (
20.     SELECT ''' + CONVERT(varchar(25), @current_date, 121) + ''' AS created_date, DB_NAME() AS
    database_name, planForceDetails.query_id,
21.     (planForceDetails.regressedPlanExecutionCount +
    planForceDetails.recommendedPlanExecutionCount)
22.     * (planForceDetails.regressedPlanCpuTimeAverage -
    planForceDetails.recommendedPlanCpuTimeAverage)/1000000 as estimated_gain,
23.     TR.reason, TR.score, JSON_VALUE(details, '$.implementationDetails.script') as script,
24.     planForceDetails.regressedPlanId, planForceDetails.recommendedPlanId,
25.     planForceDetails.regressedPlanCpuTimeAverage,
    planForceDetails.recommendedPlanCpuTimeAverage,
26.     planForceDetails.regressedPlanExecutionCount, planForceDetails.recommendedPlanExecutionCount
27. FROM sys.dm_db_tuning_recommendations AS TR
28. CROSS APPLY OPENJSON (Details, '$.planForceDetails')
29.     WITH ( [query_id] int '$.queryId',
30.           regressedPlanId int '$.regressedPlanId',
31.           recommendedPlanId int '$.recommendedPlanId',
32.           regressedPlanErrorCount int,
33.           recommendedPlanErrorCount int,
34.           regressedPlanExecutionCount int,
35.           regressedPlanCpuTimeAverage float,
36.           recommendedPlanExecutionCount int,
37.           recommendedPlanCpuTimeAverage float
38.         ) AS planForceDetails
39. LEFT JOIN sys.query_store_query AS Q ON planForceDetails.query_id = Q.query_id
40. JOIN sys.query_store_query_text AS QT ON Q.query_text_id = QT.query_text_id
41. LEFT JOIN sys.query_store_plan AS regressedQP ON planForceDetails.regressedPlanId =
    regressedQP.plan_id
42. LEFT JOIN sys.query_store_plan AS recQP ON planForceDetails.recommendedPlanId = recQP.plan_id
43. WHERE JSON_VALUE(state, '$.currentValue') = 'Active'
44. AND planForceDetails.regressedPlanCpuTimeAverage/1000.0 >= 500 --bad plan uses at least 500
    milliseconds for CPU time
45. AND
46. planForceDetails.regressedPlanCpuTimeAverage/planForceDetails.recommendedPlanCpuTimeAverage >= 10 --
    bad plan is at least 10 times slower than the good plan
47. AND planForceDetails.recommendedPlanExecutionCount >= 500 --good plan has at least 500
    executions
48. --AND planForceDetails.regressedPlanExecutionCount >= 200 --bad plan has at least 200
    executions
49. AND TR.reason < 'Number of errors in the regressed plan is greater than the recommended
    plan'
50. ) t
51. WHERE estimated_gain >= 1000
52. ORDER BY estimated_gain desc, score desc;
53.
54. EXEC sp_ineachdb @cmd, @user_only = 1, @exclude_list = @database;
55.

```


What happens next?



- <https://rb.gy/p1bme2>



Summary – Setting up Query Store

- Which defaults to change (maybe)
- ...and which to leave alone
- Trace flags
- Set up monitoring job

Summary – Tools you can use

- SSMS Built-in Reports
- Forced plans are rarely a long-term solution
- T-SQL queries
- sp_BlitzQueryStore, sp_QuickieStore

Summary – Problems you can solve

- Testing index changes
- Reducing pain from SQL Server upgrades
- Deadlocks
- Parameter sniffing

That's the end. Thank you!



desertrdb.com



jeff@desertrdb.com



[github.com\desertrdb](https://github.com/desertrdb)

